EDUCATIONAL REVIEW ARTICLE

Essential Biostatistical Methods for Epidemiologists in the Era of Artificial Intelligence and Robotics: *An Educational Review*

Olumide M. Arigbede, DrPH, MPH

Office of the Senior Vice President for Research, SUNY Downstate Health Sciences University, Brooklyn, New York, USA

ABSTRACT

The convergence of epidemiology, biostatistics, artificial intelligence (AI), and robotics represents one of the most consequential methodological shifts in the history of population health science. As data grow larger, faster, and more heterogeneous, the epidemiologist's analytical repertoire must extend well beyond classical frequency measures and significance tests. This educational review systematically characterizes the spectrum of biostatistical competencies that both practicing and emerging epidemiologists must cultivate—from foundational probability theory and regression modeling through survival analysis, longitudinal methods, causal inference, spatial epidemiology, Bayesian inference, and machine learning—and situates each within the realities imposed by AI-augmented research environments and robotic data-collection platforms. Original figures and tables accompany each major section to aid comprehension, and all equations are presented with complete symbol definitions. A structured learning roadmap, calibrated to progressive career stages, concludes the review. The central argument is definite: quantitative rigor is not merely one option among many; it is the most durable professional investment an epidemiologist can make.

Keywords: biostatistics; epidemiology; artificial intelligence; machine learning; causal inference; survival analysis; Bayesian methods; spatial epidemiology; regression modeling; robotics in health research

SECTION 1

Introduction

Epidemiology has always been a quantitative discipline. From John Snow's meticulous tabulation of cholera deaths in mid-nineteenth-century London to the genome-wide association studies of the twenty-first century, the capacity to translate observations about disease in human populations into defensible numerical conclusions has defined the craft.[1] For most of the twentieth century, the methodological core of this discipline rested on a relatively compact toolkit: measures of disease frequency, effect estimates from regression models, and probability-based tests of statistical hypotheses. That compact toolkit remains indispensable, but it is no longer sufficient.

Two forces, in particular, are reshaping what epidemiologists must know and do. The first is artificial intelligence (AI), specifically the branch concerned with machine learning, which offers methods capable of detecting patterns in data whose dimensionality, complexity, and volume would overwhelm any classical analytical strategy.[2] The second is robotics, which is rapidly generating entirely new classes of health-relevant data, from surgical performance metrics to wearable sensor streams, that demand novel analytical designs.[3] Together, these forces do not supersede traditional biostatistics; they amplify its

importance by demanding that practitioners understand when and why AI outputs are trustworthy, and when they are not.

This review is addressed equally to seasoned epidemiologists navigating a transformed methodological landscape and to graduate students and early-career scientists deciding which competencies to prioritize. It proceeds systematically through each major domain of biostatistics, defines every key concept in accessible language before providing technical detail, introduces original figures and tables to anchor conceptual understanding, and closes with a practical developmental roadmap. No prior familiarity beyond introductory statistics is assumed for the definitional material; subsequent technical sections presuppose at least one graduate-level methods course.

SECTION 2

Defining the Key Terms

Before proceeding to technical content, precision in language is warranted. Each keyword in the title of this review carries a specific meaning that shapes the argument to follow.

BIOSTATISTICAL TOOLS Mathematical and computational procedures used to collect, summarize, and draw reliable conclusions from health and biological data. Think of them as the calibrated instruments a craftsperson reaches for; selecting the wrong instrument produces a wrong answer regardless of the quality of the underlying data.	**EPIDEMIOLOGIST** A scientist who studies how, where, and why diseases and health conditions occur in human populations, and what factors increase or decrease their likelihood. Epidemiologists identify risk, inform prevention, and evaluate interventions.
ARTIFICIAL INTELLIGENCE A broad term for computer systems that perform tasks such as recognizing images, translating text, and predicting outcomes that ordinarily require human judgment. In health research, AI most often refers to statistical learning algorithms that identify patterns in large, complex datasets without being explicitly programmed for each pattern.	**ROBOTICS** The engineering and science of machines that can sense their environment, make decisions, and execute physical or computational actions with minimal human intervention. In health research, robots perform procedures, collect sensor data, and operate laboratory equipment, all of which generate rich longitudinal datasets amenable to epidemiological analysis.
MACHINE LEARNING A subset of AI in which algorithms improve their predictive performance automatically by exposure to data, rather than through explicit instruction. Most modern AI applications in health research are, at their core, machine learning systems.	**CAUSAL INFERENCE** A family of methods for estimating the effect of one variable on another in a way that can be interpreted as a cause-and-effect relationship, rather than merely a statistical association. It is the methodological bridge between observational data and experimental conclusions.

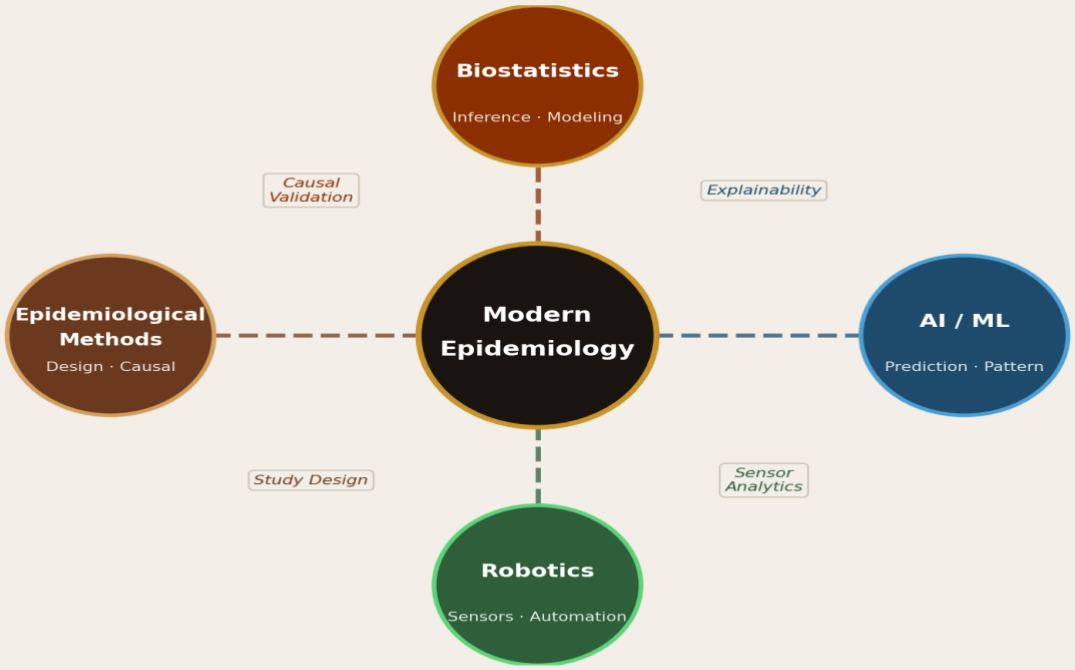

Figure 1. The four-domain nexus of modern epidemiology. Each satellite domain connects bidirectionally to the central practice through labeled bridging relationships.

Figure 1. The four-domain nexus of modern epidemiology. The central node represents the integrated practice of contemporary population health science. Each satellite represents a domain whose methods and data streams are bidirectionally linked with the central practice. Bridging labels on the connector edges name the specific analytical relationships that require cross-domain fluency. Biostatistics: provides inference and modeling; Epidemiological Methods: supplies study design and causal framing; AI/ML: contributes prediction and pattern recognition, which requires statistical explainability; Robotics: generates sensor and automation data, which require tailored analytical pipelines.

SECTION 3

Foundational Biostatistical Concepts

Every advanced analytical technique rests on a foundation of descriptive statistics, probability distributions, and inferential reasoning. Proficiency at this foundational level is not optional; it is the precondition for recognizing when more sophisticated methods are and are not warranted.

3.1 Descriptive Statistics and Measures of Central Tendency

Descriptive statistics organize and summarize raw data into comprehensible forms without inferring beyond the observed sample. The arithmetic mean, the most familiar summary measure, is defined as:

Eq. (1) $$\bar{x} = (1/n) \times \Sigma\, x_i \quad (i = 1 \text{ to } n)$$

where $\bar{x}$ is the sample mean, n is the total number of observations, and x_i is the value of the i-th observation. The Greek capital letter Σ denotes summation over all n values.

The sample variance and its square root, the standard deviation, quantify the dispersion of observations around the mean:

Eq. (2) $$s^2 = [1/(n-1)] \times \Sigma(x_i - \bar{x})^2 \;;\; s = \sqrt{s^2}$$

where s^2 is the sample variance, s is the sample standard deviation, and the denominator $n-1$ applies Bessel's correction to produce an unbiased estimate of the population variance σ^2.

The median and interquartile range (IQR) are preferable when the data distribution is markedly skewed or contains influential outliers, both of which are common in health datasets (e.g., hospital length-of-stay, health care expenditures).

3.2 Probability Distributions

Statistical inference is inseparable from probability theory. The normal distribution is the cornerstone of parametric inference and is given by the probability density function:

Eq. (3) $$f(x \mid \mu, \sigma^2) = (1 / \sigma\sqrt{2\pi}) \times \exp[-(x-\mu)^2 / (2\sigma^2)]$$

where μ is the population mean, σ^2 is the population variance, σ is the population standard deviation, $\pi \approx 3.14159$, and exp(·) denotes the natural exponential function. The distribution is symmetric about μ and fully characterized by these two parameters.

Discrete count outcomes encountered in disease surveillance, such as the number of cases per week or the number of emergency-department visits per day, are often modeled with the Poisson distribution (parameter λ, representing the expected count) or the negative binomial distribution, the latter accommodating overdispersion when $\mathrm{Var}(Y) > \mathrm{E}(Y)$.[4]

3.3 Hypothesis Testing and Confidence Intervals

A confidence interval at level (1−α) for a population mean, assuming a large sample, takes the form:

Eq. (4) $$\bar{x} \pm z(\alpha/2) \times (s / \sqrt{n})$$

where $z(\alpha/2)$ is the critical value from the standard normal distribution corresponding to a two-tailed type I error rate of α (for a 95% CI, $z_{0.025} = 1.96$), s is the sample standard deviation, and $\sqrt{n}$ is the square root of the sample size. The resulting interval is correctly interpreted as: were the study repeated many times under identical conditions, approximately (1−α)×100% of such intervals would contain the true population mean μ.

Table 1. Core descriptive and inferential biostatistical tools: plain-language descriptions, epidemiological applications, representative software commands (R and SAS), and primary analytical limitations. Software package names are italicized; functions are in monospace font.

Method	Plain-Language Description	Epidemiological Application	R / SAS Implementation	Primary Limitation
Arithmetic Mean	Average of all observed values	Describing mean age of a cohort, mean BMI in a survey sample	mean(x) / PROC MEANS	Sensitive to outliers; misleading for skewed data
Median & IQR	Middle value; spread of the middle 50% of data	Summarizing income, length of hospital stay, viral load	median(x); IQR(x) / PROC UNIVARIATE	Less amenable to mathematical manipulation than the mean
Standard Deviation	Typical distance of observations from the mean	Quantifying variability in blood pressure measurements	sd(x) / STD option in PROC MEANS	Same units as the data; not scale-free
Chi-Square Test	Tests whether two categorical variables are independent	Comparing vaccination status between cases and controls	chisq.test() / PROC FREQ CHISQ	Requires sufficient expected cell counts (≥5); not valid for sparse tables

Method	Plain-Language Description	Epidemiological Application	R / SAS Implementation	Primary Limitation
Two-Sample t-Test	Tests whether two group means differ significantly	Comparing mean cholesterol between exposed and unexposed	t.test() / PROC TTEST	Assumes approximate normality; sensitive to unequal variances
ANOVA	Generalizes the t-test to three or more groups	Comparing mean viral load across multiple treatment arms	aov(); lm() / PROC GLM	Assumes homogeneity of variance; does not identify which groups differ
Pearson Correlation	Linear association between two continuous variables	Assessing correlation between age and diastolic blood pressure	cor(x,y) / PROC CORR	Sensitive to outliers; captures only linear relationships
Spearman Correlation	Rank-based association; robust to outliers	Non-parametric association between ordinal health scores	cor(x,y, method="spearman") / PROC CORR SPEARMAN	Less powerful than Pearson when normality holds
95% Confidence Interval	A plausible range for the population parameter	Reporting uncertainty around an odds ratio or incidence rate	confint() / CLM or CLPARM options	Frequently misinterpreted as a probability statement about a fixed parameter
Mann–Whitney U-Test	Non-parametric comparison of two group distributions	Comparing days-to-recovery between two treatment groups	wilcox.test() / PROC NPAR1WAY WILCOXON	Lower statistical power than the t-test when normality holds

IQR = interquartile range; ANOVA = analysis of variance. All R base functions are available without additional packages.

SECTION 4

Regression Modeling

Regression analysis is the workhorse of epidemiological inference. By modeling the relationship between an outcome variable and one or more predictor variables — while holding others constant — regression enables effect estimation in the presence of confounding, a near-universal feature of observational health data.

4.1 Multiple Linear Regression

When the outcome is continuous and approximately normally distributed, the multiple linear regression model is expressed as:

Eq. (5) $$Y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \ldots + \beta_p X_{ip} + \varepsilon_i$$

where Y_i is the outcome for subject i; β_0 is the intercept (expected value of Y when all predictors equal zero); β_j is the partial regression coefficient for the j-th predictor X_{ij}, interpreted as the expected change in Y per one-unit increase in X_j holding all other predictors constant; $\varepsilon_i \sim N(0, \sigma^2)$ is the random error term assumed to be independently and identically distributed. Ordinary least squares (OLS) estimates $\hat{\beta} = (X'X)^{-1}X'Y$ minimize the residual sum of squares.

4.2 Logistic Regression

Binary outcomes (disease/no disease) require a model that constrains predicted values to the interval [0,1]. Logistic regression accomplishes this via the logit link function:

Eq. (6) $$log[\pi_i / (1-\pi_i)] = \beta_0 + \beta_1 X_{i1} + \ldots + \beta_p X_{ip}$$

where $\pi_i = P(Y_i = 1 \mid X_i)$ is the probability of the outcome for subject i; the left-hand side is the log-odds or logit. Exponentiating a coefficient yields the odds ratio (OR): $OR_j = e^{\beta_j}$, the multiplicative change in the odds of outcome per

one-unit increase in X_j. Coefficients are estimated by maximum likelihood via iteratively reweighted least squares (IRWLS).[5]

4.3 Poisson and Negative Binomial Regression

Count outcomes, for instance, disease cases per person-year, emergency department visits per week, among others, are modeled with Poisson regression, which uses the natural logarithm as the link function and includes an offset term to account for unequal follow-up time:

Eq. (7) $$log(\mu_i) = log(t_i) + \beta_0 + \beta_1 X_{i1} + \ldots + \beta_p X_{ip}$$

where $\mu_i = E(Y_i)$ is the expected count for subject i; $\log(t_i)$ is the offset term representing the log of the person-time at risk (its coefficient is fixed at 1); e^{β_j} is the incidence rate ratio (IRR) for a one-unit increase in X_j. When the empirical variance exceeds the mean — overdispersion — negative binomial regression introduces a dispersion parameter θ such that $Var(Y_i) = \mu_i + \mu_i^2/\theta$, accommodating the additional variance.[4] In SAS: `PROC GENMOD DIST=POISSON` or `DIST=NEGBIN`; in R: `glm(family=poisson)` or `MASS::glm.nb()`.

Practical note for the modern epidemiologist: When AI systems are applied to the same data used in a regression analysis: for example, a neural network trained to predict the same binary outcome as a logistic regression, the regression model often remains the scientifically preferred output for causal or inferential questions because its coefficients carry interpretable meaning, uncertainty estimates, and assumptions that can be explicitly evaluated. The AI model may outperform it on predictive accuracy in held-out data, but prediction and causal estimation are distinct objectives that require distinct methods.[6]

SECTION 5

Survival Analysis

Survival analysis (also called time-to-event analysis) is the analytical framework for outcomes defined by the time elapsed until a specified event occurs, such as death, disease onset, hospital discharge, or treatment failure. Its distinguishing methodological feature is the formal accommodation of censoring: the situation in which the event has not yet occurred for some subjects by the end of the study period, or when subjects are lost to follow-up. Ignoring censoring by analyzing only observed events produces systematically biased estimates.[7]

5.1 The Kaplan–Meier Estimator

The Kaplan–Meier (KM) estimator provides a non-parametric estimate of the survival function S(t) = P(T > t), the probability that an individual survives past time t:

Eq. (8) $$\hat{S}(t) = \Pi[t_j \leq t]\,(1 - d_j/n_j)$$

where the product is taken over all event times $t_j \leq t$; d_j is the number of events (e.g., deaths) occurring at time t_j; and n_j is the number of individuals at risk (i.e., not yet experienced the event and not yet censored) immediately before time t_j.[8] The resulting step function is visualized as a staircase that descends at each event time. Between two groups, the log-rank test evaluates whether their KM curves are statistically compatible with arising from the same underlying survival distribution.

5.2 The Cox Proportional Hazards Model

When the effect of covariates on survival time is of primary interest, the Cox proportional hazards (CPH) model is the standard approach:

Eq. (9) $$h(t \mid X) = h_0(t) \times \exp(\beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_p X_p)$$

where $h(t \mid X)$ is the hazard function, defined as the instantaneous rate of the event at time t for an individual with covariate vector X; $h_0(t)$ is the unspecified baseline hazard function (the hazard when all covariates equal zero); and e^β_j is the hazard ratio (HR) for a one-unit increase in X_j, holding other covariates constant.[9] Coefficients are estimated by maximizing the partial likelihood, which conditions on the observed event order without requiring $h_0(t)$ to be estimated. The proportional hazards assumption, which states that HR is constant over time, must be formally assessed using scaled Schoenfeld residuals and, if violated, addressed through time-varying coefficients or stratification. In R: `survival::coxph()`; in SAS: `PROC PHREG`.

5.3 Competing Risks

In many clinical and epidemiological settings, subjects may experience one of several mutually exclusive events, for example, death from cancer versus death from cardiovascular disease, and the occurrence of one event precludes observation of the other. The standard KM estimator overestimates the cause-specific cumulative incidence in such settings. The Fine–Gray sub-distribution hazard model provides the appropriate framework by modeling the hazard of the cumulative incidence function (CIF) for each event type directly, keeping subjects who experienced a competing event in the risk set.[10] In R: `cmprsk::crr()` or `tidycmprsk::crr()`; in SAS: `PROC PHREG` with the `EVENTCODE` option.

SECTION 6

Longitudinal Data Analysis and Time Series Methods

Longitudinal studies, in which the same subjects are measured repeatedly over time, are indispensable for establishing temporal sequences of exposure and outcome, which is a prerequisite for causal argumentation. Their analysis requires methods that explicitly model the correlation structure of repeated measurements within individuals.

6.1 Linear Mixed-Effects Models

The linear mixed-effects model (LMM) partitions variability into population-level fixed effects and subject-specific random effects:

Eq. (10) $$Y_i = X_i\beta + Z_i b_i + \varepsilon_i$$

where Y_i is the vector of repeated outcomes for subject i; X_i is the fixed-effects design matrix with coefficient vector β; Z_i is the random-effects design matrix; $b_i \sim N(0, G)$ is the vector of subject-specific random effects with covariance matrix G; and $\varepsilon_i \sim N(0, R_i)$ is the residual error vector with covariance matrix R_i.[11] Parameters are typically estimated via restricted maximum likelihood (REML). In R: `lme4::lmer()`; in SAS: `PROC MIXED` or `PROC GLIMMIX`.

6.2 Generalized Estimating Equations

When the primary inferential target is the population-average effect rather than the subject-specific effect, generalized estimating equations (GEE) offer a robust semiparametric alternative. GEE extends generalized linear models to clustered or longitudinal data by specifying a working correlation structure — exchangeable, autoregressive, or unstructured — and producing consistent estimates of β even when that working correlation structure is misspecified, provided the mean model is correct, and the sandwich variance estimator is used.[12] In R: `geepack::geeglm()`; in SAS: `PROC GENMOD REPEATED`.

6.3 ARIMA Time Series Modeling

Syndromic surveillance data, usually defined as daily or weekly counts of emergency department visits, influenza-like illness, or overdose events, among many others, constitute time series, and their analysis requires models that explicitly account for temporal autocorrelation. The autoregressive integrated moving average (ARIMA) model of order (p, d, q) is expressed using the backshift operator B (where $B^k Y_t = Y_{t-k}$):

Eq. (11) $$\varphi_p(B) \times (1-B)^d \times Y_t = \theta q(B) \times \varepsilon_t$$

where $\phi p(B) = 1 - \phi_1 B - \phi_2 B^2 - \ldots - \phi p B^p$ is the autoregressive (AR) polynomial of order p; $(1-B)^d$ is the differencing operator applied d times to achieve stationarity; $\theta q(B) = 1 + \theta_1 B + \theta_2 B^2 + \ldots + \theta q B^q$ is the moving-average (MA) polynomial of order q; and $\varepsilon t \sim N(0, \sigma^2 \varepsilon)$ is white noise. The seasonal generalization, SARIMA(p,d,q)(P,D,Q)s, adds seasonal AR, differencing, and MA terms at lag s. In public health applications, interrupted time series (ITS) analysis embeds a policy or intervention indicator within an ARIMA model, enabling evaluation of intervention effects while controlling for pre-existing temporal trends.[13]

SECTION 7

Causal Inference

Statistical association, however precisely estimated, does not establish causation. Epidemiology's central mission, which is the identification of causes of disease to inform prevention, demands methods that can support causal conclusions from observational data, an objective requiring frameworks that go substantially beyond standard regression.[14]

7.1 Directed Acyclic Graphs

A directed acyclic graph (DAG) is a visual representation of the investigator's causal assumptions about the relationships among variables in a study. Nodes represent variables; directed edges (arrows) represent hypothesized causal pathways. DAGs explicitly identify which variables are confounders (common causes of exposure and outcome), mediators (variables on the causal pathway), and colliders (variables with two or more causes), and they identify the minimal adjustment set required to estimate an unconfounded causal effect using the backdoor criterion.[15] Tools: DAGitty (dagitty.net); R package `dagitty`.

7.2 Propensity Score Methods

The propensity score, introduced by Rosenbaum and Rubin,[16] is the conditional probability of treatment assignment given observed baseline covariates: $e(X) = P(A=1 \mid X)$. Under the assumption of no unmeasured confounding (strong ignorability), conditioning on the propensity score — through matching, inverse probability weighting (IPW), or stratification — produces a covariate balance that approximates the conditions of a randomized experiment. The stabilized IPW estimator weights each subject by:

Eq. (12) $$w_i = P(A_i)/P(A_i|X_i) = A_i/\hat{e}(X_i) + (1-A_i)/(1-\hat{e}(X_i))$$

where $A_i \in \{0,1\}$ is the treatment indicator for subject i; $\hat{e}(X_i)$ is the estimated propensity score; and $P(A_i)$ is the marginal probability of treatment (often estimated as the sample proportion). The weighted pseudo-population has covariate distributions balanced between treatment groups. In R: `MatchIt` for matching; `WeightIt` for weighting.

7.3 Instrumental Variables and Mendelian Randomization

When unmeasured confounding cannot be ruled out by design or propensity-score methods, instrumental variable (IV) analysis offers an alternative. A valid instrument Z must satisfy three conditions: (1) Z is associated with the exposure A; (2) Z is independent of all unmeasured confounders; and (3) Z affects the outcome Y only through A (the exclusion restriction). Mendelian randomization exploits the random assortment of alleles at conception as a natural instrument for modifiable exposures (e.g., serum lipids, adiposity), leveraging the fact that genotype is determined before the development of confounders.[17] Two-stage least squares (2SLS) is the standard estimator; in R: `ivreg::ivreg()` and `MendelianRandomization::mr_ivw()`.

SECTION 8

Bayesian Biostatistics

Bayesian inference provides a unified framework for incorporating prior knowledge with observed data to produce posterior probability distributions over unknown parameters. Its philosophical foundation (that probability represents a degree of belief that is updated as evidence accumulates) contrasts with the frequentist interpretation underlying confidence intervals and *p*-values, but the two frameworks converge numerically in large samples with diffuse priors.

8.1 Bayes' Theorem

Eq. (13) $$\pi(\theta|Y) = L(Y|\theta)\times\pi(\theta) \,/\, \int L(Y|\theta)\pi(\theta)d\theta \;\; \alpha \;\; L(Y|\theta)\times\pi(\theta)$$

where π(θ | Y) is the posterior distribution of the parameter vector θ given observed data Y; L(Y | θ) is the likelihood function; π(θ) is the prior distribution encoding pre-data knowledge; and the denominator is the marginal likelihood (normalizing constant), which ensures the posterior integrates to one. In practice, Markov chain Monte Carlo (MCMC) algorithms, including Gibbs sampling, Metropolis–Hastings, or Hamiltonian Monte Carlo as implemented in Stan, are used to draw samples from the posterior when analytical integration is intractable.[18]

Bayesian methods are particularly valuable in epidemiology when sample sizes are small (as in rare-disease outbreak investigations), hierarchical (multilevel) structures are present, and adaptive clinical trial designs continuously update evidence thresholds. In R: `rstan`, `brms`; in Python: `PyMC`.

SECTION 9

Machine Learning in Epidemiological Research

Machine learning (ML) methods are not a replacement for traditional biostatistical models; they are a complementary set of tools optimized for different objectives. Classical regression prioritizes parameter interpretability and uncertainty quantification. ML prioritizes predictive accuracy across complex, high-dimensional feature spaces. The epidemiologist's task is to deploy the right tool for the right question.[2,6,29]

9.1 Regularized Regression: LASSO

The least absolute shrinkage and selection operator (LASSO) modifies ordinary least squares by adding an L_1 penalty to the objective function, simultaneously performing variable selection and coefficient shrinkage:

Eq. (14) $$\hat{\beta}LASSO = argmin \{ \Sigma(Y_i - X_i'\beta)^2 + \lambda \, \Sigma|\beta_i| \}$$

where $\lambda \geq 0$ is the regularization hyperparameter (selected by cross-validation); the L_1 penalty $\lambda\Sigma|\beta j|$ sets coefficients for irrelevant predictors exactly to zero, producing a sparse model. When $\lambda = 0$, the solution reduces to OLS. The elastic net combines L_1 (LASSO) and L_2 (ridge) penalties, trading off sparsity for stability when predictors are highly correlated. In R: `glmnet::glmnet()`; in Python: `sklearn.linear_model.Lasso`.[19]

9.2 Random Forests

Random forests construct an ensemble of decision trees, each trained on a bootstrap resample of the data and a random subset of predictors at each split, and aggregate their predictions:

Eq. (15) $$\hat{f}(x) = (1/B) \times \Sigma[b=1 \text{ to } B]\ T_\beta(x;\ \Theta_\beta)$$

where $\hat{f}(x)$ is the ensemble prediction for covariate vector x; B is the number of trees; $T_\beta(x; \Theta b)$ is the b-th decision tree trained on a bootstrap sample with random feature subset Θb at each split; and the average is taken over all trees (for continuous outcomes) or replaced by a majority vote (for binary outcomes).[20] Variable importance, quantified as the mean decrease in node impurity (Gini importance) or the permutation-based mean decrease in accuracy, provides a scale-free ranking of predictors. In R: `ranger::ranger()` or `randomForest::randomForest()`.

9.3 Neural Networks and Deep Learning

A feedforward neural network computes a prediction by composing nonlinear transformations across layers of artificial neurons:

Eq. (16) $$a(l) = g(l)[\ W(l) \times a(l-1) + b(l)\ ]\ ,\ \ l = 1, 2, \ldots, L$$

where $a^{(l)}$ is the activation vector at layer l; $W^{(l)}$ is the weight matrix and $b^{(l)}$ the bias vector for layer l; $g^{(l)}(\cdot)$ is the activation function—commonly ReLU(z) = max(0,z) for hidden layers and the sigmoid or softmax for output layers; and $a^{(0)} = x$ is the input vector. Parameters $\{W^{(l)}, b^{(l)}\}$ are estimated by backpropagating gradients with stochastic gradient descent or adaptive optimizers (e.g., Adam, RMSProp). Deep learning architectures, including convolutional, recurrent, and transformer-based, extend this framework to image, sequence, and text data, respectively.[21]

Table 2. Comparative overview of classical biostatistical and machine learning methods: analytical paradigm, orientation, interpretability, key assumptions, representative epidemiological application, and primary software implementation. This table is intended to guide method selection rather than to declare superiority of any one approach.

Method	Paradigm	Primary Orientation	Interpretability	Key Assumption(s)	Epidemiological Application	Software
Linear Regression	Classical	Inference	High	Linearity, normality, homoscedasticity, independence	Estimating mean blood pressure change per BMI unit	R: lm(); SAS: PROC REG
Logistic Regression	Classical	Inference	High	Logit-linear relationship, no separation, independence	Estimating OR for disease risk factors in case-control study	R: glm(binomial); SAS: PROC LOGISTIC
Cox PH Model	Classical	Inference	High	Proportional hazards, non-informative censoring	Estimating HR for cancer survival by treatment group	R: survival::coxph(); SAS: PROC PHREG
Poisson / NegBin	Classical	Inference	High	Log-linear mean; Poisson requires equidispersion	Modeling weekly flu case counts in syndromic surveillance	R: glm(poisson); MASS::glm.nb()
Linear Mixed Model	Classical	Inference	High	Normality of random effects; correct covariance specification	Analyzing repeated CD4 count measurements in an HIV cohort	R: lme4::lmer(); SAS: PROC MIXED

Method	Paradigm	Primary Orientation	Interpretability	Key Assumption(s)	Epidemiological Application	Software
LASSO / Elastic Net	ML / Regularized	Prediction + selection	Moderate	Sparse true signal; pre-specified λ via CV	Variable selection in high-dimensional genomic data	R: glmnet; Python: sklearn
Random Forest	ML / Ensemble	Prediction	Low–Moderate	None (non-parametric); prone to overfitting without tuning	Predicting 30-day hospital readmission from EHR features	R: ranger; Python: sklearn
Gradient Boosting (XGBoost)	ML / Ensemble	Prediction	Low–Moderate	Requires careful hyperparameter tuning; computationally intensive	Mortality risk scoring in intensive care unit data	R: xgboost; Python: xgboost
Convolutional Neural Net	Deep Learning	Prediction	Very Low	Large labeled training data; GPU resources	Automated detection of lesions on radiological images	Python: TensorFlow, PyTorch
K-Means Clustering	Unsupervised ML	Exploration	Moderate	Spherical clusters; pre-specified k; sensitive to outliers	Identifying patient phenotypes in electronic health records	R: kmeans(); Python: sklearn
Principal Component Analysis	Unsupervised ML	Dimensionality reduction	Moderate	Linear structure; variance as a proxy for information	Adjusting for population stratification in genome-wide association studies	R: prcomp(); Python: sklearn

OR = odds ratio; HR = hazard ratio; PH = proportional hazards; NegBin = negative binomial; EHR = electronic health record; CV = cross-validation; GPU = graphics processing unit; LASSO = least absolute shrinkage and selection operator.

SECTION 10

Spatial Epidemiology and Geographic Information Systems

Geographic clustering of disease is a cardinal feature of infectious disease epidemiology, environmental health, and health disparity research. Spatial epidemiology applies statistical methods to characterize where disease rates are elevated, whether that elevation exceeds chance, and what local risk factors may explain observed patterns.[22]

10.1 Moran's I: Measuring Spatial Autocorrelation

Eq. (17) $$I = [n / \Sigma_i\Sigma_j w_{ij}] \times [\Sigma_i\Sigma_j w_{ij}(x_i-\bar{x})(x_j-\bar{x})] / [\Sigma_i(x_i-\bar{x})^2]$$

where n is the number of spatial units (e.g., counties, census tracts); x_i and x_j are the outcome values at locations i and j; $\bar{x}$ is the overall mean; and w_{ij} are elements of the spatial weights matrix W, typically set to 1 if locations i and j share a boundary (queen contiguity) and 0 otherwise. Moran's I ∈ [−1, +1], where values significantly above zero indicate positive spatial autocorrelation (similar values cluster together), values near zero indicate spatial randomness, and negative values indicate spatial dispersion. In R: `spdep::moran.test()`.[22]

Beyond global spatial autocorrelation, the Kulldorff spatial scan statistic identifies specific geographic clusters of elevated or reduced disease risk by evaluating likelihood ratios for all possible circular windows across the study area, with inference based on Monte Carlo simulation.[23] Tools: SaTScan (freely available from the National Cancer Institute); R package `SpatialEpi`. The Bayesian Besag–York–Mollié (BYM) model further accommodates spatially structured and unstructured random effects in disease mapping within an integrated nested Laplace approximation (INLA) framework; in R: `INLA::inla()`.

SECTION 11

Systematic Review and Meta-Analysis

A single study, however well-designed, yields an estimate subject to sampling variability. Meta-analysis combines evidence across studies to produce a pooled effect estimate with substantially greater precision and the capacity to characterize heterogeneity, a variation in true effects across populations, interventions, and settings.[24]

11.1 The Random-Effects Model

The DerSimonian–Laird random-effects model treats each study as estimating a study-specific true effect θk drawn from a distribution of true effects:

Eq. (18) $$\hat{\theta}RE = \Sigma\, wk^* \hat{\theta}k \,/\, \Sigma\, wk^* \;;\; wk^* = 1/(vk + \hat{\tau}^2)$$

where K is the number of studies; $\hat{\theta}k$ is the effect estimate from study k (e.g., log OR, log RR, mean difference); $vk = \hat{V}ar(\hat{\theta}k)$ is the within-study variance; $\hat{\tau}^2$ is the estimated between-study variance (heterogeneity), estimated by the method of moments (DerSimonian–Laird) or restricted maximum likelihood (REML); and $w^*k = 1/(vk + \hat{\tau}^2)$ is the inverse-variance weight. When $\hat{\tau}^2 = 0$, the model reduces to the fixed-effect estimator. Heterogeneity is quantified by $I^2 = \hat{\tau}^2/(\hat{\tau}^2 + \tilde{v}) \times 100\%$, where $\tilde{v}$ is a typical within-study variance; $I^2 \geq 75\%$ signals substantial heterogeneity.[24] In R: `metafor::rma()`;`meta::metagen()`.

Publication bias, the tendency for studies with statistically significant or large-effect results to be published at higher rates than null-result studies, threatens the validity of meta-analytic conclusions. Funnel plot asymmetry, Egger's linear regression test, and Begg's rank correlation test are the standard diagnostic tools, with trim-and-fill and Copas selection model sensitivity analyses available to quantify potential bias.[24]

SECTION 12

AI, Robotics, and the Practicing Epidemiologist

The interface between AI systems and epidemiological practice is not merely technical; it is also ethical and scientific. The epidemiologist's role in an AI-augmented environment is to evaluate, not simply adopt, algorithmic outputs.

12.1 AI in Disease Surveillance

Natural language processing (NLP) and social media analytics now complement traditional surveillance by enabling near-real-time extraction of disease signals from unstructured text: emergency department chief complaints, clinical notes, news reports, and internet search queries. Platforms such as HealthMap and the World Health Organization's Event Information Site (EIS) deploy such systems globally.[25] However, AI surveillance outputs must be evaluated using the same metrics applied to traditional case definitions: sensitivity, specificity, positive predictive value, and timeliness. An algorithm with high sensitivity that performs poorly during novel pathogen emergence, to be precise, the scenario in which early detection matters most, offers illusory safety.

12.2 Epidemiological Analysis of Robotic Surgical Data

Robotic surgical platforms (e.g., the da Vinci Surgical System) generate timestamped, high-resolution operative metrics, such as instrument trajectories, force measurements, and task completion times, that

constitute a new class of longitudinal health data.[3,28] Epidemiological analyses of such data address learning curves, complication clustering by surgeon or institution, and comparative effectiveness of robotic versus conventional approaches. Appropriate methods include multilevel models (to account for surgeon-level clustering), competing-risks analysis (to compare complications with successful completion), and ITS designs (to evaluate technology adoption effects at the institutional level).

12.3 Algorithmic Bias and Epistemic Responsibility

A landmark analysis by Obermeyer and colleagues[26] demonstrated that a widely deployed commercial algorithm used to identify high-risk patients for care management, which used healthcare costs as a proxy for health need, systematically underestimated the true health burden of Black patients because structural inequities cause them to incur lower costs at equivalent health levels. This finding, which carried direct implications for health disparities, would never have surfaced without rigorous epidemiological scrutiny. Algorithmic fairness evaluation, even including stratifying model performance metrics by race, sex, age, and socioeconomic status, is not a peripheral data science concern; it is a core epidemiological competency.

12.4 Explainable AI

Methods for explainable AI (XAI), not limited to SHAP (SHapley Additive exPlanations) values, LIME (Local Interpretable Model-agnostic Explanations), and attention visualization in transformer models, provide post-hoc explanations of individual predictions generated by otherwise opaque algorithms.[27] For clinical decision support and public health policy applications, where accountability and scientific defensibility are non-negotiable, epidemiologists should demand explanations alongside predictions and critically assess whether those explanations are biologically or contextually plausible.

SECTION 13

A Structured Learning Roadmap for Epidemiologists

No practitioner masters the entire spectrum described in this review simultaneously. The following four-stage roadmap is calibrated to progressive career development, from the beginning of graduate training through independent research practice. Each stage builds on the prior and should be revisited periodically as methodological literature evolves.

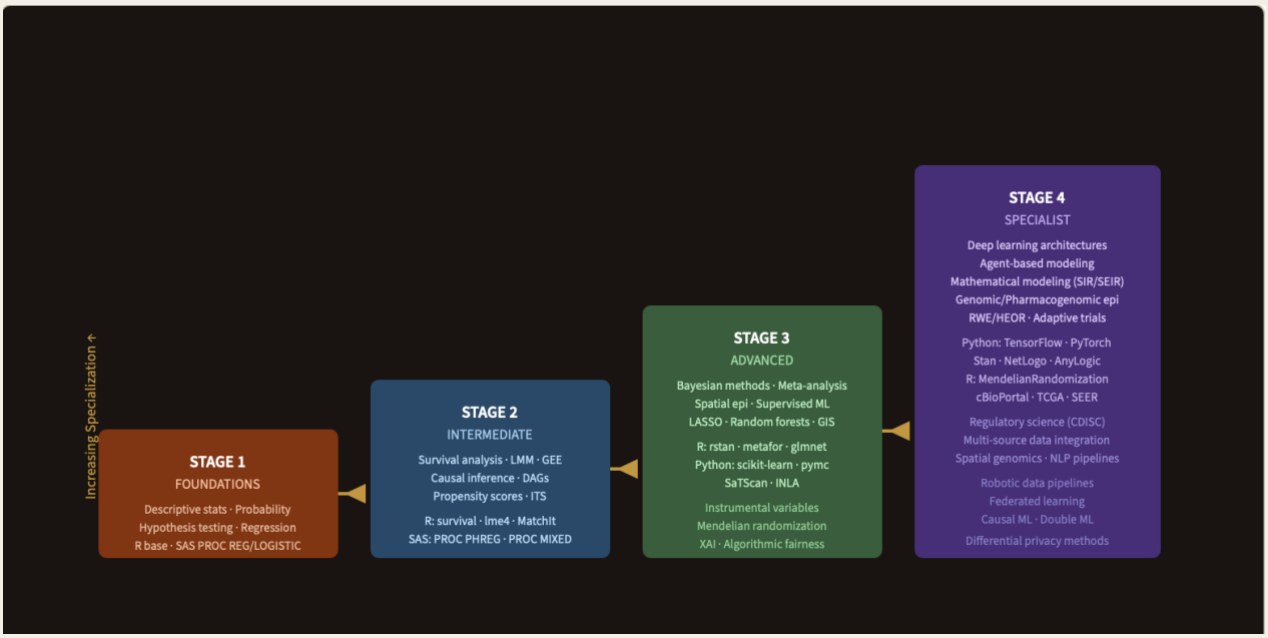

Figure 2. A four-stage progressive learning roadmap for epidemiologists. Stages advance from left (foundational) to right (specialist), corresponding broadly to pre-master's, mid-master's, doctoral, and post-doctoral or mid-career practice levels, respectively. Each stage lists representative methods, software tools, and platforms. Practitioners at any career stage should revisit earlier stages when encountering a new substantive area. Color coding: red = Stage 1 (Foundations); blue = Stage 2 (Intermediate); green = Stage 3 (Advanced); purple = Stage 4 (Specialist). RWE = real-world evidence; HEOR = health economics and outcomes research; TCGA = The Cancer Genome Atlas; SEER = Surveillance, Epidemiology, and End Results.

SECTION 14

Conclusion

The analytical landscape confronting epidemiologists in the current era is more demanding, consequential, and intellectually rewarding than at any prior point in the history of the discipline. The methods reviewed here, from the arithmetic mean to Hamiltonian Monte Carlo, from the Kaplan–Meier estimator to convolutional neural networks, do not compete with one another; they constitute an integrated analytical vocabulary, each term appropriate to a different type of question, data structure, or inferential objective.

Artificial intelligence and robotics do not diminish the importance of biostatistical foundations; they intensify it. A practitioner who cannot distinguish a predictive model from a causal one, who cannot recognize when an AI algorithm's training data violates the assumptions required for valid inference in a new population, or who cannot identify when a time series requires differencing before modeling is a practitioner who will eventually produce a consequential error. The sophistication of the algorithmic wrapper around such an error does not mitigate its harm.

Conversely, the epidemiologist who commands the spectrum of methods described in this review is uniquely positioned to contribute to public health in this era: capable of designing rigorous studies, extracting credible causal estimates from complex observational data, deploying and critically

evaluating AI systems, and communicating quantitative uncertainty with precision to policymakers and the public alike. Cultivating that command is among the most consequential professional investments an epidemiologist can make.

REFERENCES

1. *Snow J. On the Mode of Communication of Cholera.* 2nd ed. London: John Churchill; 1855.

2. Topol EJ. High-performance medicine: the convergence of human and artificial intelligence. *Nat Med.* 2019;25(1):44–56. doi:10.1038/s41591-018-0300-7

3. Leal Ghezzi T, Campos Corleta O. 30 years of robotic surgery. *World J Surg.* 2016;40(10):2550–2557. doi:10.1007/s00268-016-3543-9

4. Hilbe JM. *Negative Binomial Regression.* 2nd ed. Cambridge: Cambridge University Press; 2011.

5. Hosmer DW, Lemeshow S, Sturdivant RX. *Applied Logistic Regression.* 3rd ed. Hoboken, NJ: Wiley; 2013.

6. Shmueli G. To explain or to predict? *Stat Sci.* 2010;25(3):289–310. doi:10.1214/10-STS330

7. Kleinbaum DG, Klein M. *Survival Analysis: A Self-Learning Text.* 3rd ed. New York: Springer; 2012.

8. Kaplan EL, Meier P. Nonparametric estimation from incomplete observations. *J Am Stat Assoc.* 1958;53(282):457–481. doi:10.1080/01621459.1958.10501452

9. Cox DR. Regression models and life-tables. *J R Stat Soc Ser B.* 1972;34(2):187–220. doi:10.1111/j.2517-6161.1972.tb00899.x

10. Fine JP, Gray RJ. A proportional hazards model for the subdistribution of a competing risk. *J Am Stat Assoc.* 1999;94(446):496–509. doi:10.1080/01621459.1999.10474144

11. Laird NM, Ware JH. Random-effects models for longitudinal data. *Biometrics.* 1982;38(4):963–974. doi:10.2307/2529876

12. Liang K-Y, Zeger SL. Longitudinal data analysis using generalized linear models. *Biometrika.* 1986;73(1):13–22. doi:10.1093/biomet/73.1.13

13. Wagner AK, Soumerai SB, Zhang F, Ross-Degnan D. Segmented regression analysis of interrupted time series studies in medication use research. *J Clin Pharm Ther.* 2002;27(4):299–309. doi:10.1046/j.1365-2710.2002.00430.x

14. Hernán MA, Robins JM. *Causal Inference: What If.* Boca Raton, FL: Chapman & Hall/CRC; 2020.

15. Greenland S, Pearl J, Robins JM. Causal diagrams for epidemiologic research. *Epidemiology.* 1999;10(1):37–48. doi:10.1097/00001648-199901000-00008

16. Rosenbaum PR, Rubin DB. The central role of the propensity score in observational studies for causal effects. *Biometrika.* 1983;70(1):41–55. doi:10.1093/biomet/70.1.41

17. Davey Smith G, Ebrahim S. 'Mendelian randomization': can genetic epidemiology contribute to understanding environmental determinants of disease? *Int J Epidemiol.* 2003;32(1):1–22. doi:10.1093/ije/dyg070

18. Gelman A, Carlin JB, Stern HS, Dunson DB, Vehtari A, Rubin DB. *Bayesian Data Analysis.* 3rd ed. Boca Raton, FL: CRC Press; 2013.

19. Tibshirani R. Regression shrinkage and selection via the lasso. *J R Stat Soc Ser B.* 1996;58(1):267–288. doi:10.1111/j.2517-6161.1996.tb02080.x

20. Breiman L. Random forests. *Mach Learn.* 2001;45(1):5–32. doi:10.1023/A:1010933404324

21. LeCun Y, Bengio Y, Hinton G. Deep learning. *Nature.* 2015;521(7553):436–444. doi:10.1038/nature14539

22. Waller LA, Gotway CA. *Applied Spatial Statistics for Public Health Data.* Hoboken, NJ: Wiley; 2004.

23. Kulldorff M. A spatial scan statistic. *Commun Stat Theory Methods.* 1997;26(6):1481–1496. doi:10.1080/03610929708831995

24. Borenstein M, Hedges LV, Higgins JPT, Rothstein HR. *Introduction to Meta-Analysis.* Chichester: Wiley; 2009.

25. Brownstein JS, Freifeld CC, Madoff LC. Digital disease detection — harnessing the web for public health surveillance. *N Engl J Med.* 2009;360(21):2153–2157. doi:10.1056/NEJMp0900702

26. Obermeyer Z, Powers B, Vogeli C, Mullainathan S. Dissecting racial bias in an algorithm used to manage the health of populations. *Science.* 2019;366(6464):447–453. doi:10.1126/science.aax2342

27. Lundberg SM, Lee S-I. A unified approach to interpreting model predictions. *Adv Neural Inf Process Syst.* 2017;30.

28. Arigbede O, Amusa T, Buxbaum SG. Exploring the use of artificial intelligence and robotics in prostate cancer management. *Cureus.* 2023;15(9):e46021. doi:10.7759/cureus.46021

29. Obermeyer Z, Emanuel EJ. Predicting the future — big data, machine learning, and clinical medicine. *N Engl J Med.* 2016;375(13):1216–1219. doi:10.1056/NEJMp1606181

© 2026 Olumide M. Arigbede, DrPH, MPH · SUNY Downstate Health Sciences University, Brooklyn, NY · Educational Review · Open Access